

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv(r"C:\Users\ARUTHRA D\Downloads\archive (1)\Titanic-Dataset.csv")
print(df)
```

	PassengerId	Survived	Pclass	\					
0	1	0	3						
1	2	1	1						
2	3	1	3						
3	4	1	1						
4	5	0	3						
..						
886	887	0	2						
887	888	1	1						
888	889	0	3						
889	890	1	1						
890	891	0	3						
					Name	Sex	Age	SibSp	\
0				Braund, Mr. Owen Harris	male	22.0	1		
1	Cummings, Mrs. John Bradley (Florence Briggs Th...			female	38.0	1			
2	Heikkinen, Miss. Laina			female	26.0	0			
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)			female	35.0	1			
4	Allen, Mr. William Henry			male	35.0	0			
..		
886	Montvila, Rev. Juozas			male	27.0	0			
887	Graham, Miss. Margaret Edith			female	19.0	0			
888	Johnston, Miss. Catherine Helen "Carrie"			female	NaN	1			
889	Behr, Mr. Karl Howell			male	26.0	0			
890	Dooley, Mr. Patrick			male	32.0	0			
	Parch	Ticket	Fare	Cabin	Embarked				
0	0	A/5 21171	7.2500	NaN	S				
1	0	PC 17599	71.2833	C85	C				
2	0	STON/O2. 3101282	7.9250	NaN	S				
3	0	113803	53.1000	C123	S				
4	0	373450	8.0500	NaN	S				
..				
886	0	211536	13.0000	NaN	S				
887	0	112053	30.0000	B42	S				
888	2	W./C. 6607	23.4500	NaN	S				
889	0	111369	30.0000	C148	C				
890	0	370376	7.7500	NaN	Q				

Out[5]: (891, 12)

```
In [6]: df.describe()
```

Out[6]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [7]: # Check the columns
```

```
print(df.columns)
```

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
      dtype='object')
```

```
In [8]: # If the column is named 'Survived'
```


```
print(df['Survived'].value_counts())
```

```
Survived  
0      549  
1      342  
Name: count, dtype: int64
```

```
In [21]: # Plotting with seaborn (sns)
```

```
import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x=df['Survived'], hue=df['Pclass'])
plt.show()
```



```
In [1]: # data cleaning
```

```
In [9]: f.drop('PassengerId', inplace=True, axis=1)
df['Cabin'].count()/df.shape[0]
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[9], line 1
----> 1 f.drop('PassengerId', inplace=True, axis=1)
      2 df['Cabin'].count()/df.shape[0]

NameError: name 'f' is not defined
```

```
In [10]: df.drop('PassengerId', inplace=True, axis=1)
df['Cabin'].count()/df.shape[0]
```

```
Out[10]: 0.22895622895622897
```

```
In [11]: df.drop('Cabin',inplace=True,axis=1)#because there is more than 70% empty values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    int64
1   Pclass      891 non-null    int64
2   Name        891 non-null    object
3   Sex         891 non-null    object
4   Age         714 non-null    float64
5   SibSp       891 non-null    int64
6   Parch       891 non-null    int64
7   Ticket      891 non-null    object
8   Fare        891 non-null    float64
9   Embarked    889 non-null    object
dtypes: float64(2), int64(4), object(4)
memory usage: 69.7+ KB
```

```
In [12]: df['Age'].count()/df.shape[0]
0.8013468013468014
df.dropna(subset=['Age', 'Embarked'],inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 712 entries, 0 to 890
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    712 non-null    int64
1   Pclass      712 non-null    int64
2   Name        712 non-null    object
3   Sex         712 non-null    object
4   Age         712 non-null    float64
5   SibSp       712 non-null    int64
6   Parch       712 non-null    int64
7   Ticket      712 non-null    object
8   Fare        712 non-null    float64
9   Embarked    712 non-null    object
dtypes: float64(2), int64(4), object(4)
memory usage: 61.2+ KB
```


```
In [13]: df['Title']=df['Name'].str.extract(r'([A-Za-z]+\.)',expand=False)
df.drop('Name',inplace=True,axis=1)
pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)
df
```

333	0	3	male	16.00	2	0	345764	18.0000	S	Mr.
336	0	1	male	29.00	1	0	113776	66.6000	S	Mr.
337	1	1	female	41.00	0	0	16966	134.5000	C	Miss.
338	1	3	male	45.00	0	0	7598	8.0500	S	Mr.
339	0	1	male	45.00	0	0	113784	35.5000	S	Mr.
340	1	2	male	2.00	1	1	230080	26.0000	S	Master.
341	1	1	female	24.00	3	2	19950	263.0000	S	Miss.
342	0	2	male	28.00	0	0	248740	13.0000	S	Mr.
343	0	2	male	25.00	0	0	244361	13.0000	S	Mr.
344	0	2	male	36.00	0	0	229236	13.0000	S	Mr.
345	1	2	female	24.00	0	0	248733	13.0000	S	Miss.
346	1	2	female	40.00	0	0	31418	13.0000	S	Miss.
348	1	3	male	3.00	1	1	C.A. 37671	15.9000	S	Master.


```
In [14]: df['Survived'].value_counts()
```

```
Out[14]: Survived
0      424
1      288
Name: count, dtype: int64
```

```
In [15]: sns.countplot(data=df,x='Survived')  
plt.ylabel('Number Of Survived')  
plt.title('Survival Rate')  
plt.grid()  
plt.show()
```



```
In [16]: sns.countplot(x='Sex',data=df)
plt.ylabel('Number Of Gender')
plt.title('Gender Rate')
plt.grid()
plt.show()
```




```
In [17]: df['Sex'].value_counts()
```

```
Out[17]: Sex
male      453
female    259
Name: count, dtype: int64
```


```
In [18]: df['Pclass'].value_counts()
```

```
Out[18]: Pclass
3      355
1      184
2      173
Name: count, dtype: int64
```


```
In [19]: sns.countplot(x='Pclass',data=df)
plt.ylabel('Number Of Classes')
plt.title('Classes')
plt.show()
```




```
In [20]: sns.histplot(df['Age'],bins=25)
plt.yticks(np.arange(0,100,10))
plt.xticks(np.arange(0,90,5))
plt.title('Age Distribution',fontsize=14)
plt.xlabel('No Of People')
plt.show()
```




```
In [21]: sns.countplot(x='SibSp',data=df)  
plt.show()
```



```
In [22]: sns.countplot(x='Parch',data=df)
```

```
Out[22]: <Axes: xlabel='Parch', ylabel='count'>
```



```
In [23]: print('Siblings:\n',df['SibSp'].value_counts())
print('Parents : \n',df['Parch'].value_counts())
```

```
Siblings:
 SibSp
0     469
1     183
2      25
4      18
3      12
5       5
Name: count, dtype: int64
Parents :
 Parch
0     519
1     110
2      68
5       5
3       5
4       4
6       1
Name: count, dtype: int64
```

```
In [24]: df.min()
```

```
Out[24]: Survived      0
Pclass      1
Sex      female
Age      0.42
SibSp      0
Parch      0
Ticket    110152
Fare      0.0
Embarked    C
Title     Capt.
dtype: object
```

```
In [26]: df.max()
```

```
Out[26]: Survived      1
Pclass      3
Sex      male
Age     80.0
SibSp      5
Parch      6
Ticket    WE/P 5735
Fare    512.3292
Embarked    S
Title     Sir.
dtype: object
```


```
In [27]: df.mode()
```

```
Out[27]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Title
0	0	3	male	24.0	0	0	347082	13.0	S	Mr.

```
In [28]: ind=df[df['Age']>30].index
df_1=df.drop(ind,axis=0)
sns.catplot(x='Survived',hue='Sex',kind='count',data=df_1)
plt.show()
```

C:\Users\ARUTHRA D\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)




```
In [29]: df.groupby('Survived')['Sex'].value_counts()
```

```
Out[29]: Survived  Sex
0             male    360
           female     64
1             female   195
           male      93
Name: count, dtype: int64
```

```
In [30]: sns.catplot(x='Survived',hue='Sex',col='Pclass',data=df,kind='count')
plt.yticks(np.arange(0,250,20))
plt.show()
```

C:\Users\ARUTHRA D\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)




```
In [31]: df.groupby(['Survived','Pclass'])['Sex'].value_counts()
```

```
Out[31]: Survived  Pclass  Sex
0          1      male      61
          1    female       3
          2      male      84
          2    female       6
          3      male     215
          3    female      55
1          1    female      80
          1      male      40
          2    female      68
          2      male      15
          3    female      47
          3      male      38
Name: count, dtype: int64
```

```
In [32]: sns.catplot(x='Survived',hue='Sex',col='Embarked',data=df,kind='count')
plt.yticks(np.arange(0,330,30))
plt.show()
```

C:\Users\ARUTHRA D\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)



```
In [33]: df.groupby(['Survived', 'Embarked'])['Sex'].value_counts()
```

```
Out[33]:
```

Survived	Embarked	Sex	
0	C	male	45
		female	6
	Q	male	15
		female	5
	S	male	300
		female	53
1	C	female	55
		male	24
	Q	female	7
		male	1
	S	female	133
		male	68

Name: count, dtype: int64


```
In [34]: sns.catplot(x='Survived',hue='Sex',col='Embarked',row='Pclass',data=df,kind='count')
plt.yticks(np.arange(0,190,20))
plt.tight_layout()
plt.show()
```

C:\Users\ARUTHRA D\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

self._figure.tight_layout(*args, **kwargs)

C:\Users\ARUTHRA D\AppData\Local\Temp\ipykernel_15688\1850565211.py:3: UserWarning: The figure layout has changed to tight

plt.tight_layout()




```
In [35]: df.groupby(['Survived', 'Embarked', 'Pclass'])['Sex'].value_counts()
```

```
Out[35]:
```

Survived	Embarked	Pclass	Sex	
0	C	1	male	20
			female	1
		2	male	7
			female	5
		3	male	18
			female	5
	Q	1	male	1
			female	1
		2	male	13
			female	5
		3	male	40
			female	2
1	C	1	male	76
			female	6
			male	184
			female	45
		2	male	37
			female	16
		3	male	7
			female	1
		3	female	11
			male	7
	Q	1	female	1
			female	1
		2	female	5
			male	1
		3	female	42
			male	24
	S	1	female	60
			male	14
			female	31
			male	30
		2	female	60
			male	14
		3	female	31
			male	30

Name: count, dtype: int64


```
In [37]: sns.boxplot(df['Age'])  
plt.show()
```



```
In [38]: def boundaries(data,col,dis):  
    Q1=data[col].quantile(0.25)  
    Q3=data[col].quantile(0.75)  
    IQR=Q3-Q1  
    low=Q1-(IQR*dis)  
    upper=Q3+(IQR*dis)  
    return low,upper  
lower,upper=boundaries(df,'Age',1.5)  
print('Lower Range : ',lower,' Upper Range : ',upper)
```

Lower Range : -7.0 Upper Range : 65.0

```
In [39]: not_out=(df['Age']<upper)&(df['Age']>lower)  
df['Age'][~not_out].count()
```

Out[39]: 11

```
In [40]: df=df[not_out]
df['Agebin']=pd.cut(df['Age'],5,labels=['a','b','c','d','e'],include_lowest=True)
sns.countplot(x='Agebin',data=df)
plt.show()
```


C:\Users\ARUTHRA D\AppData\Local\Temp\ipykernel_15688\1455345314.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['Agebin']=pd.cut(df['Age'],5,labels=['a','b','c','d','e'],include_lowest=True)
```



```
In [48]: # Step 1: Import necessary Libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [45]: # Step 2: Load the dataset
file_path = r"C:\Users\ARUTHRA D\Downloads\archive (1)\Titanic-Dataset.csv"
titanic_data = pd.read_csv(file_path)
```

```

In [46]: # Step 3: Data Preprocessing
# Handling missing values
titanic_data['Age'].fillna(titanic_data['Age'].median(), inplace=True)
titanic_data['Embarked'].fillna(titanic_data['Embarked'].mode()[0], inplace=True)
titanic_data.drop(columns=['Cabin'], inplace=True)

In [49]: # Encoding categorical variables
label_encoder = LabelEncoder()
titanic_data['Sex'] = label_encoder.fit_transform(titanic_data['Sex'])
titanic_data['Embarked'] = label_encoder.fit_transform(titanic_data['Embarked'])

In [50]: # Selecting features and target variable
features = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']
X = titanic_data[features]
y = titanic_data['Survived']

In [51]: # Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scaling the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

In [52]: # Step 4: Model Building
model = LogisticRegression()
model.fit(X_train, y_train)

# Making predictions
y_pred = model.predict(X_test)

# Step 5: Model Evaluation
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

In [53]: print(f"Accuracy: {accuracy}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)

```

```

Accuracy: 0.8044692737430168
Confusion Matrix:
[[90 15]
 [20 54]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.82	0.86	0.84	105
1	0.78	0.73	0.76	74
accuracy			0.80	179
macro avg	0.80	0.79	0.80	179
weighted avg	0.80	0.80	0.80	179

```
In [54]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [56]: # importing dataset
```

```
df=pd.read_csv(r"C:\Users\ARUTHRA D\Downloads\archive (1)\Titanic-Dataset.csv")
print(df)
```

143	144	0	3
144	145	0	2
145	146	0	2
146	147	1	3
147	148	0	3
148	149	0	2
149	150	0	2
150	151	0	2
151	152	1	1
152	153	0	3
153	154	0	3
154	155	0	3
155	156	0	1
156	157	1	3
157	158	0	3
158	159	0	3
159	160	0	3
160	161	0	3
161	162	1	2
162	163	0	3

```
In [57]: # Load the dataset
titanic_data = pd.read_csv(r"C:\Users\ARUTHRA D\Downloads\archive (1)\Titanic-Dataset.csv")

# Display the first few rows of the dataset
print(titanic_data.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
In [58]: # Drop columns that are not needed for this analysis
titanic_data = titanic_data.drop(columns=['Name', 'Ticket', 'Cabin'])

# Encode categorical variables
label_encoder = LabelEncoder()
titanic_data['Sex'] = label_encoder.fit_transform(titanic_data['Sex'])
titanic_data['Embarked'] = titanic_data['Embarked'].fillna('S')
titanic_data['Embarked'] = label_encoder.fit_transform(titanic_data['Embarked'])

# Handle missing values
imputer = SimpleImputer(strategy='mean')
titanic_data['Age'] = imputer.fit_transform(titanic_data[['Age']])

# Display the preprocessed data
print(titanic_data.head())
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	1	0	3	1	22.0	1	0	7.2500	2
1	2	1	1	0	38.0	1	0	71.2833	0
2	3	1	3	0	26.0	0	0	7.9250	2
3	4	1	1	0	35.0	1	0	53.1000	2
4	5	0	3	1	35.0	0	0	8.0500	2

```
In [59]: # Define features and target variable
X = titanic_data.drop(columns=['Survived'])
y = titanic_data['Survived']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [60]: # Initialize and train the Random Forest classifier
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)
```

```
In [61]: # Print classification report and confusion matrix
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.83      0.88      0.85        105
     1       0.81      0.74      0.77         74


 accuracy          0.82          179
 macro avg         0.82      0.81      0.81          179
 weighted avg      0.82      0.82      0.82          179
```


```
Confusion Matrix:
[[92 13]
 [19 55]]
```

```
In [62]: # Plot feature importances
feature_importances = pd.DataFrame(model.feature_importances_, index=X.columns, columns=

plt.figure(figsize=(10, 6))
sns.barplot(x=feature_importances.index, y=feature_importances['importance'])
plt.title('Feature Importances')
plt.xlabel('Features')
plt.ylabel('Importance')
plt.xticks(rotation=45)
plt.show()

# Plot confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```





In []: